

# Decision Trees

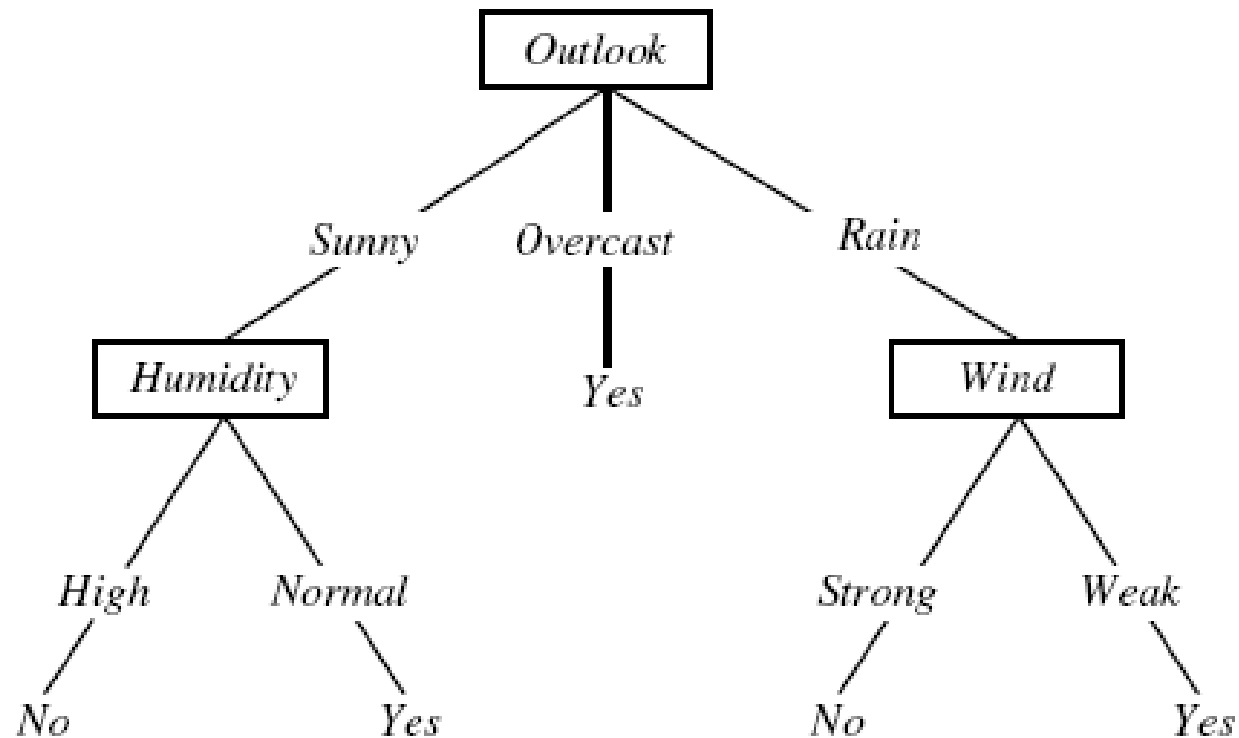
# Review

- We introduced three styles of algorithms for linear classifiers:
  - 1) Perceptron – learn classifier direction
  - 2) Logistic Regression – learn  $P(Y | X)$
  - 3) Linear Discriminant Analysis – learn  $P(X, Y)$
- Linear separability
  - A data set is linearly separable if there exists a linear hyperplane that separates positive examples from negative examples

# Not linearly separable data

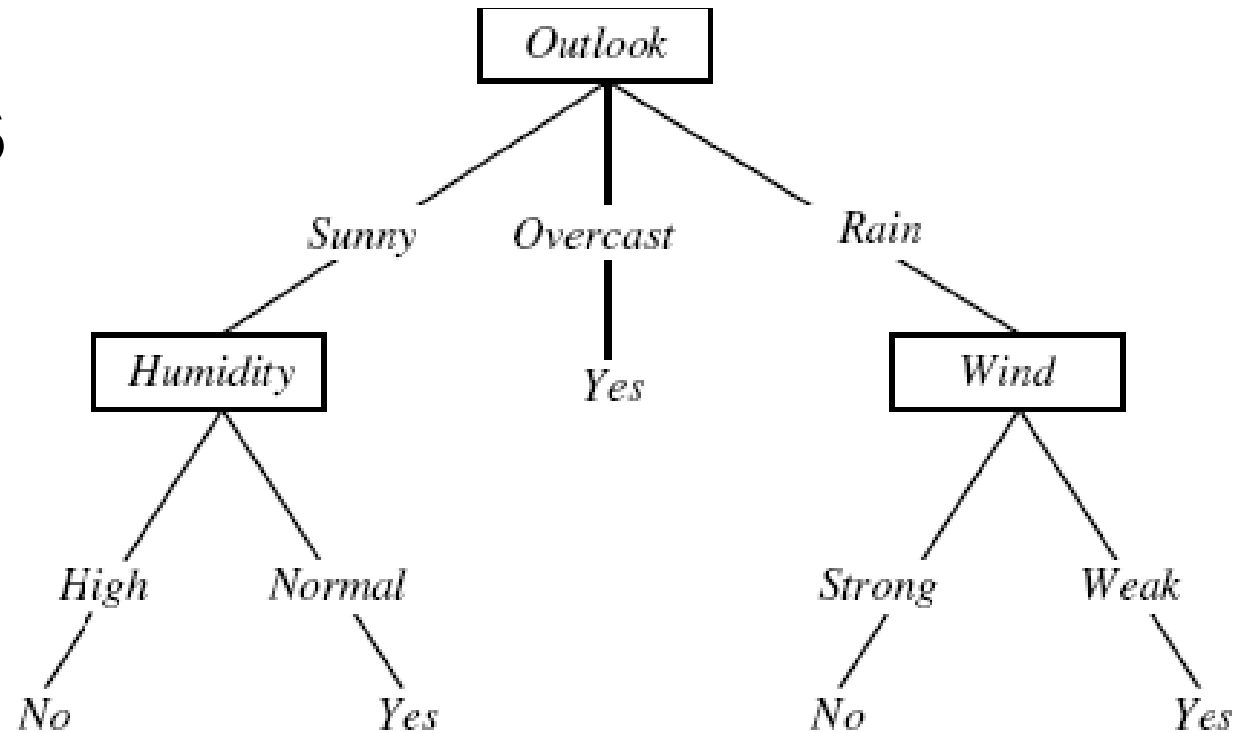
- Some data sets are not linearly separable!
- Option 1
  - Use non-linear features, e.g., polynomial basis functions
  - Learn linear classifiers in the non-linear feature space
  - Will discuss more later
- Option 2
  - Use non-linear classifiers (decision trees, neural networks, nearest neighbors etc.)

# Decision Tree for Playing Tennis



Prediction is done by sending the example down the tree till a class assignment is reached



# Definitions



- **Internal nodes**
  - Each test a feature
  - Branch according to feature values
  - Discrete features – branching is naturally defined
  - Continuous features – branching by comparing to a threshold
- **Leaf nodes**
  - Each assign a classification

# Decision Tree Decision Boundaries



- Decision Trees divide the feature space into axis-parallel rectangles and label each rectangle with one of the  $K$  classes



# Hypothesis Space of Decision Trees

- Decision trees provide a very popular and efficient hypothesis space
  - Deal with both **Discrete** and **Continuous** features
  - **Variable size**: as the # of nodes (or depth) of tree increases, the hypothesis space grows
    - Depth 1 (“decision stump”) can represent any Boolean function of one feature
    - Depth 2: Any Boolean function of two features and some Boolean functions involving three features:
    - In general, can represent any Boolean functions

# Decision Trees Can Represent Any Boolean Function



- If a target Boolean function has  $n$  inputs, there always exists a decision tree representing that target function.
- However, in the worst case, exponentially many nodes will be needed (why?)
  - $2^n$  possible inputs to the function
  - In the worst case, we need to use one leaf node to represent each possible input



# Learning Decision Trees

- Goal: Find a decision tree  $h$  that achieves minimum misclassification errors on the training data
- A trivial solution: just create a decision tree with one path from root to leaf for each training example
  - Bug: Such a tree would just memorize the training data. It would not generalize to new data points
- Solution 2: Find the smallest tree  $h$  that minimizes error
  - Bug: This is NP-Hard

# Top-down Induction of Decision Trees

There are different ways to construct trees from data. We will focus on the top-down, greedy search approach:

Basic idea:




1. Choose the **best** feature  $a^*$  for the root of the tree.
2. Separate training set  $\mathbf{S}$  into subsets  $\{\mathbf{S}_1, \mathbf{S}_2, \dots, \mathbf{S}_k\}$  where each subset  $\mathbf{S}_i$  contains examples having the same value for  $a^*$ .
3. Recursively apply the algorithm on each new subset until all examples have the same class label.

# Choosing Feature Based on Classification Error


- Perform 1-step look-ahead search and choose the attribute that gives the lowest error rate on the training data

$x_1$	$x_2$	$x_3$	$y$
0	0	0	1
0	0	1	0
0	1	0	1
0	1	1	1
1	0	0	0
1	0	1	1
1	1	0	0
1	1	1	0

Training examples



Unfortunately, this measure does not always work well, because it does not detect cases where we are making “progress” toward a good tree



# Entropy

- Let  $X$  be a random variable with the following probability distribution

$P(X = 0)$	$P(X = 1)$
0.2	0.8


- The entropy of  $\mathbf{X}$ , denoted  $H(\mathbf{X})$ , is defined as

$$H(X) = - \sum_x P_X(x) \log_2 P_X(x)$$

- Entropy** measures the uncertainty of a random variable
- The larger the entropy, the more uncertain we are about the value of  $X$
- If  $P(X=0)=0$  (or 1), there is no uncertainty about the value of  $X$ , entropy = 0
- If  $P(X=0)=P(X=1)=0.5$ , the uncertainty is maximized, entropy = 1

# Entropy

- Entropy is a concave function downward



# More About Entropy

- Joint Entropy

$$H(X, Y) = - \sum_x \sum_y P(X = x, Y = y) \log P(X = x, Y = y)$$

- Conditional Entropy is defined as

$$\begin{aligned} H(Y | X) &= \sum_x P(X = x) H(Y | X = x) \\ &= - \sum_x P(X = x) \sum_y P(Y = y | X = x) \log P(Y = y | X = x) \end{aligned}$$

– The average surprise of Y when we know the value of X

- Entropy is additive


$$H(X, Y) = H(X) + H(Y | X)$$

# Mutual Information

- The mutual information between two random variables  $X$  and  $Y$  is defined as:

$$I(X, Y) = H(Y) - H(Y | X)$$

- the amount of information we learn about  $Y$  by knowing the value of  $X$  (and vice versa – it is symmetric).
- Consider the class  $Y$  of each training example and the value of feature  $X_1$  to be random variables. The mutual information quantifies how much  $X_1$  tells us about  $Y$ .





# Choosing the Best Feature


- Choose the feature  $X_j$  that has the highest mutual information with  $Y$  - often referred to as the **information gain** criterion

$$\begin{aligned}\arg \max_j I(X_j; Y) &= \arg \max_j H(Y) - H(Y | X_j) \\ &= \arg \min_j H(Y | X_j)\end{aligned}$$

- Define  $\tilde{J}(j)$  to be the expected remaining uncertainty about  $y$  after testing  $x_j$

$$\tilde{J}(j) = H(Y | X_j) = \sum_x P(X_j = x) H(Y | X_j = x)$$

# Choosing the Best Feature: A General View



$$\text{Benefit of split} = U(S) - [P_L * U(S_L) + P_R * U(S_R)]$$


Expected Remaining Uncertainty (Impurity)

Measures of Uncertainty	
Error	$\min\{p, 1 - p\}$
Entropy	$-p \log p - (1 - p) \log 1 - p$
Gini Index	$2p(1 - p)$

# Multi-nomial Features

- Multiple discrete values
  - Method 1: Construct multi-way split
    - Information Gain will tend to prefer multi-way split
    - To avoid this bias, we rescale the information gain:

**Correction:** the correct objective should be  $\arg\max I(X_j, Y)/H(X_j)$


$$\arg \min_j \frac{H(Y | X_j)}{H(X_j)} = \arg \min_j \frac{\sum_x P(X_j = x) H(Y | X_j = x)}{-\sum_x P(X_j = x) \log P(X_j = x)}$$

- Method 2: Test for one value versus all of the others
- Method 3: Group the values into two disjoint sets and test one set against the other

# Continuous Features

- Test against a threshold
- How to compute the best threshold  $\theta_j$  for  $X_j$ ?
  - Sort the examples according to  $X_j$ .
  - Move the threshold  $\theta$  from the smallest to the largest value
  - Select  $\theta$  that gives the best information gain
  - Only need to compute information gain when class label changes

# Continuous Features

- Information gain for  $\theta = 1.2$  is 0.2294

Y	0	0	1	0	1	1	0	1	1
X <sub>j</sub>	0.2	0.4	0.7	1.1	1.3	1.7	1.9	2.4	2.9

$n_{0,L} = 3$	$n_{0,R} = 1$
$n_{1,L} = 1$	$n_{1,R} = 4$

- Information gain only needs to be computer when class label changes

Y	0	0	1	0	1	1	0	1	1
X <sub>j</sub>	0.2	0.4	0.7	1.1	1.3	1.7	1.9	2.4	2.9

# Top-down Induction of Decision Trees


There are different ways to construct trees from data. We will focus on the top-down, greedy search approach:

Basic idea:

1. Choose the **best** feature  $a^*$  for the root of the tree.
2. Separate training set  $\mathbf{S}$  into subsets  $\{\mathbf{S}_1, \mathbf{S}_2, \dots, \mathbf{S}_k\}$  where each subset  $\mathbf{S}_i$  contains examples having the same value for  $a^*$ .
3. Recursively apply the algorithm on each new subset until all examples have the same class label.


# Over-fitting

- Decision tree has a very flexible hypothesis space
- As the nodes increase, we can represent arbitrarily complex decision boundaries – training set error is always zero
- This can lead to over-fitting



Possibly just noise, but the tree is grown larger to capture these examples

# Over-fitting






# Avoid Overfitting

- Early stop
  - Stop growing the tree when data split does not differ significantly different from random split
- Post pruning
  - Separate training data into **training set** and **validating set**
  - Evaluate impact on validation set when pruning each possible node
  - Greedily remove the one that most improve the validation set performance

# Effect of Pruning



# Decision Tree Summary

- DT - one of the most popular machine learning tools
  - Easy to understand
  - Easy to implement
  - Easy to use
  - Computationally cheap
- Information gain to select features (ID3, C4.5 ...)
- DT over-fits
  - Training error can always reach zero
- To avoid overfitting:
  - Early stopping
  - Pruning