

Dimension Reduction

CS534

Why dimension reduction?

- High dimensionality – large number of features
 - E.g., documents represented by thousands of words, millions of bigrams
 - Images represented by thousands of pixels
- Redundant and irrelevant features (not all words are relevant for classifying/clustering documents)
- Difficult to interpret and visualize
- Curse of dimensionality

Extract Latent Linear Features

- Linearly project n -d data onto a k -d space
 - e.g., project space of 10^4 words into 3-dimensions
- There are infinitely many k -d subspaces that we can project the data into, which one should we choose
- This depends on the task at hand
 - If supervised learning, we would like to maximize the separation among classes: Linear discriminant analysis (LDA)
 - If unsupervised, we would like to retain as much data variance as possible: principal component analysis (PCA)


LDA: linear discriminant analysis

- Also named Fisher Discriminant Analysis
- It can be viewed as
 - *a dimension reduction* method
 - a generative classifier ($p(x|y)$): Gaussian with distinct μ for each class but shared Σ
- We will now look at its dimension reduction interpretation


Intuition

- Find a project direction so that the separation between classes is maximized
- In other words, we are looking for a projection that best discriminates different classes

Classes mixed



Better Separation



Objectives of LDA

- One way to measure separation is to look at the class means

$$\mathbf{m}_1 = \frac{1}{N_1} \sum_{\mathbf{x} \in c_1} \mathbf{x} \quad \mathbf{m}_2 = \frac{1}{N_2} \sum_{\mathbf{x} \in c_2} \mathbf{x}$$

Original
means

$$m'_1 = \frac{1}{N_1} \sum_{\mathbf{x} \in c_1} \mathbf{w}^T \mathbf{x} \quad m'_2 = \frac{1}{N_2} \sum_{\mathbf{x} \in c_2} \mathbf{w}^T \mathbf{x}$$

Projected
means

$$|m'_1 - m'_2|^2 = |\mathbf{w}^T \mathbf{m}_1 - \mathbf{w}^T \mathbf{m}_2|^2$$

We want the distance between the projected means to be as large as possible

Objectives of LDA

- We further want the data points from the same class to be as close as possible
- This can be measured by the class ***scatter*** (*variance within the class*)

$$s_i^2 = \sum_{x \in c_i} (\mathbf{w}^T \mathbf{x} - m'_i)^2$$

Total within class scatter for projected class i

$$s_1^2 + s_2^2$$

Total within class scatter

Combining the two sides

- There are a number of different ways to combine these two sides of the objective
- LDA seeks to optimize the following objective:

$$\arg \max_w \frac{|m'_1 - m'_2|^2}{s_1^2 + s_2^2}$$

Diagram illustrating the components of the LDA objective function:

- The numerator $|m'_1 - m'_2|^2$ is circled and points to the equation:

$$|m'_1 - m'_2|^2 = (w^T m_1 - w^T m_2)^2$$

$$= w^T (m_1 - m_2)(m_1 - m_2)^T w$$

$$= w^T S_B w$$
- The denominator $s_1^2 + s_2^2$ is circled and points to the equation:

$$s_1^2 + s_2^2 = w^T (S_1 + S_2) w$$

$$= w^T S_W w$$

$$s_1^2 = \sum_{x \in C_1} (w^T x - w^T m_1)^2 = \sum_x w^T (x - m_1)(x - m_1)^T w$$

$$= w^T \left(\sum_x (x - m_1)(x - m_1)^T \right) w = w^T S_1 w$$

The LDA Objective

$$J(\mathbf{w}) = \frac{\mathbf{w}^T S_B \mathbf{w}}{\mathbf{w}^T S_w \mathbf{w}}$$

$$S_i = \sum_{x \in C_i} (x - m_i)(x - m_i)^T$$

$$S_B = (m_1 - m_2)(m_1 - m_2)^T$$

the between class scatter matrix

$$S_w = S_1 + S_2$$

the total within class scatter matrix, where

$$S_i = \sum_{x \in C_i} (x - m_i)(x - m_i)^T$$

- The above objective is known as generalized Reyleigh quotient, and it's easy to show a w that maximizes $J(w)$ must satisfy $S_B w = \lambda S_w w$
- Noticing that $S_B w = (m_1 - m_2)(m_1 - m_2)^T w$ always take the direction of $m_1 - m_2$

Scalar
- Ignoring the scalars, this leads to:



$$(m_1 - m_2) = S_w w$$

$$w = S_w^{-1}(m_1 - m_2)$$

LDA for two classes

$$\mathbf{w} = S_w^{-1}(\mathbf{m}_1 - \mathbf{m}_2)$$

- Projecting data onto one dimension that maximizes the ratio of between-class scatter and total within-class scatter



LDA for Multi-Class

$$J(w) = \frac{w^T S_B w}{w^T S_W w}$$

- Objective remains the same, with slightly different definition for between-class scatter:


$$S_B = \frac{1}{k} \sum_{i=1}^k (m_i - m)(m_i - m)^T$$

m is the overall mean

- Solution: $k-1$ eigenvectors of $S_W^{-1} S_B$

Unsupervised Dimension Reduction


- Consider data without class labels
- Try to find a more compact representation of the data



$3d \Rightarrow 2d$



- Assume that the high dimensional data actually resides in a inherent low-dimensional space
- Additional dimensions are just random noise
- Goal is to recover these inherent dimensions and discard noise dimensions

Geometric picture of principal components (PCs)



Goal: to account for the variation in the data in as few dimensions as possible


Geometric picture of principal components (PCs)



- The 1st PC is the projection direction that maximizes the variance of the projected data
- The 2nd PC is the projection direction that is orthogonal to the 1st PC and maximizes the variance


Conceptual Algorithm

- Find a line such that when the data is projected onto that line, it has the maximum variance




Conceptual Algorithm

- Find a new line, orthogonal to the first, that has maximum projected variance:



Repeat Until d Lines

- The projected position of a point on these lines gives the coordinates in the m -dimensional reduced space



Computing this set of lines is achieved by eigen-decomposition of the covariance matrix S

PCA: variance maximization

- Given n data points: x_1, \dots, x_n
- Consider a linear projection specified by v
- The projection of x onto v is $z = v^T x$
- The variance of the projected data is
$$\text{var}(z) = \text{var}(v^T x v) = v^T \text{var}(x) v = v^T S v$$
- The 1st PC maximizes the variance subject to the constraint $v^T v = 1$


Maximizing Variance

- Maximize $v^T S v$, subject to $v^T v = 1$
- Lagrange:

$$v^T S v - \lambda(v^T v - 1)$$
$$\frac{\partial}{\partial v} = 0 \Rightarrow S v = \lambda v$$

- Thus v is the eigen-vector of S with eigen-value λ
- Sample variance of the projected data:
$$v^T S v = \lambda v^T v = \lambda$$
- The eigen-values = the amount of variance captured by each eigen-vector
 - 1st PC = The first eigen-vector
 - 2nd PC = the second eigen-vector
 - ...

Alternative view: Minimizing Reconstruction Error




We can view this as low-d approximation of the original data:

$$x^1 \approx x_0 + z^1 u; \quad x^2 \approx x_0 + z^2 u$$

The goal is to minimize the difference between the original data and the approximation.

Dimension Reduction Using PCA

- Calculate the covariance matrix of the data S
- Calculate the eigen-vectors/eigen-values of S
- Rank the eigen-values in decreasing order
- Select eigen-vectors that retain a fixed percentage of the variance, (e.g., 80%, the smallest d such that $\frac{\sum_{i=1}^d \lambda_i}{\sum_i \lambda_i} \geq 80\%$)




You might lose some info. But if the eigen-values are small, not much is lost.

Example: Face Recognition

- An typical image of size 256 x 128 is described by $n = 256 \times 128 = 32768$ dimensions
- Each face image lies somewhere in this high-dimensional space
- Images of faces are generally similar in overall configuration, thus
 - They cannot be randomly distributed in this space
 - We should be able to describe them in a much low-dimensional space

PCA for Face Images: Eigenfaces

- Database of 128 carefully-aligned faces.
- Here are the mean and the first 15 eigenvectors.
- Each eigenvector can be shown as an image
- These images are face-like, thus called eigenface



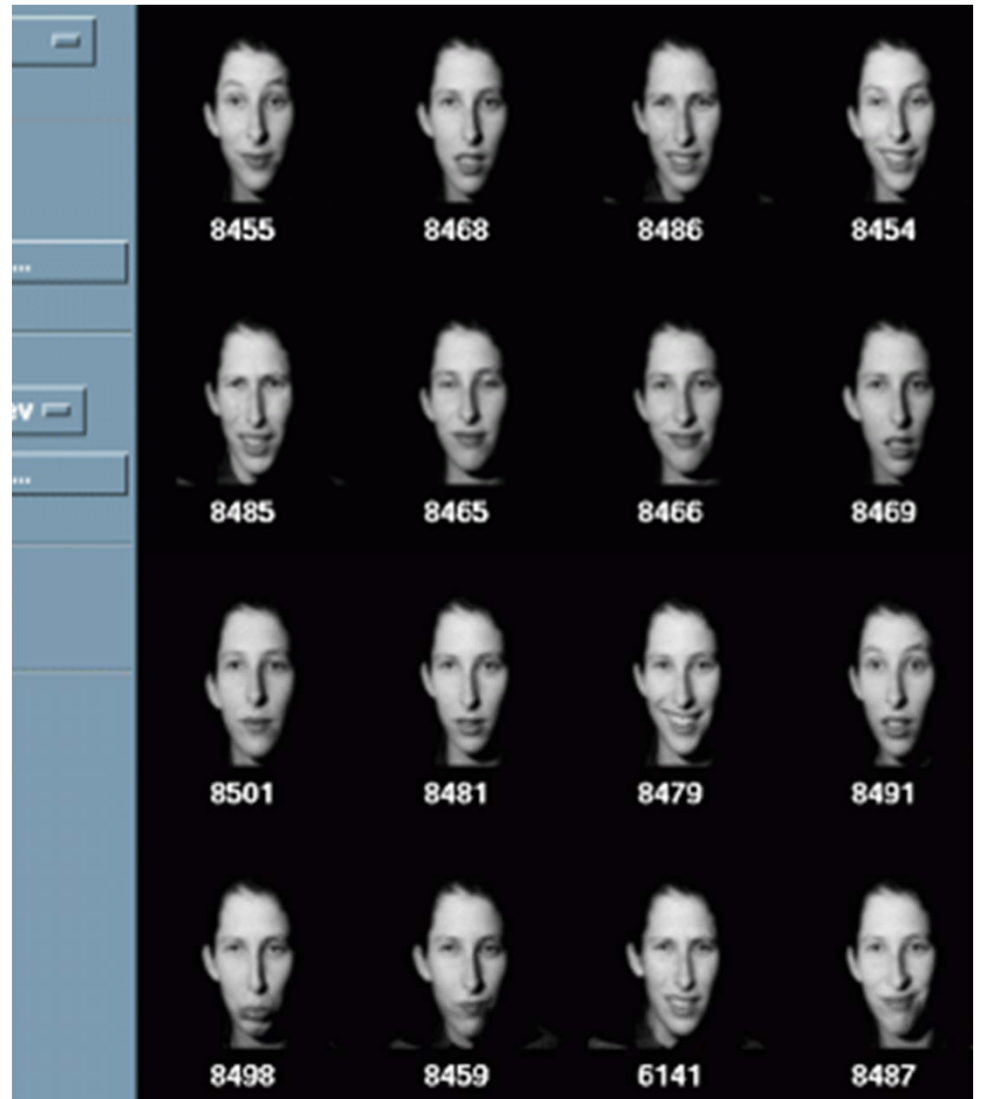
Face Recognition in Eigenface space

(Turk and Pentland 1991)

- Nearest Neighbor classifier in the eigenface space
- Training set always contains 16 face images of 16 people, all taken under the same conditions of lighting, head orientation, and image size
- Accuracy:
 - variation in lighting: 96%
 - variation in orientation: 85%
 - variation in image size: 64%

Face Image Retrieval

- Left-top image is the query image
- Return 15 nearest neighbor in the eigenface space
- Able to find the same person despite
 - different expressions
 - variations such as glasses




PCA: A Useful Preprocessing Step

- Helps to reduce the computational complexity
- Helps supervised learning
 - Reduced dimension \Rightarrow simpler hypothesis space
 - Smaller VC dimension \Rightarrow less over-fitting
- PCA can also be seen as noise reduction
- Fails when data consists of multiple separate clusters


Practical Issue: Scaling Up

- Covariance of the image data is BIG!
 - size of $\Sigma = 32768 \times 32768$
 - finding eigenvector of such a matrix is slow.
- SVD comes to rescue!
 - Can be used to compute principal components
 - Efficient implementations available, e.g., Matlab svd

Singular Value Decomposition: $X=USV^T$



Singular Value Decomposition: $X=USV^T$




SVD for PCA

- Create centered data matrix X
- Solve SVD: $X = USV^T$
- Columns of V are the eigenvectors of Σ sorted from largest to smallest eigenvalues – select the first k columns as our principal components

Nonlinear Dimension Reduction

Nonlinear Methods

- Data often lies on or near a nonlinear low-dimensional curve
- We call such low dimension structure manifolds




Swiss roll data

ISOMAP: Isometric Feature Mapping

(Tenenbaum et al. 2000)

- A nonlinear method for dimensionality reduction
- Preserves the global, nonlinear geometry of the data by preserving the geodesic distances
- Geodesic: originally geodesic means the shortest route between two points on the surface of the manifold



ISOMAP

- Two steps
 1. Approximate the geodesic distance between every pair of points in the data
 - The manifold is locally linear
 - Euclidean distance works well for points that are close enough
 - For the points that are far apart, their geodesic distance can be approximated by summing up local Euclidean distances
 2. Find a Euclidean mapping of the data that preserves the geodesic distance

Geodesic Distance

- Construct a graph by
 - Connecting i and j if
 - $d(i, j) < \varepsilon$ (ε -isomap) or
 - i is one of j 's k nearest neighbors (k -isomap)
 - Set the edge weight equal $d(i, j)$ – Euclidean distance
- Compute the Geodesic distance between any two points as the ***shortest path distance***

Compute the Low-Dimensional Mapping

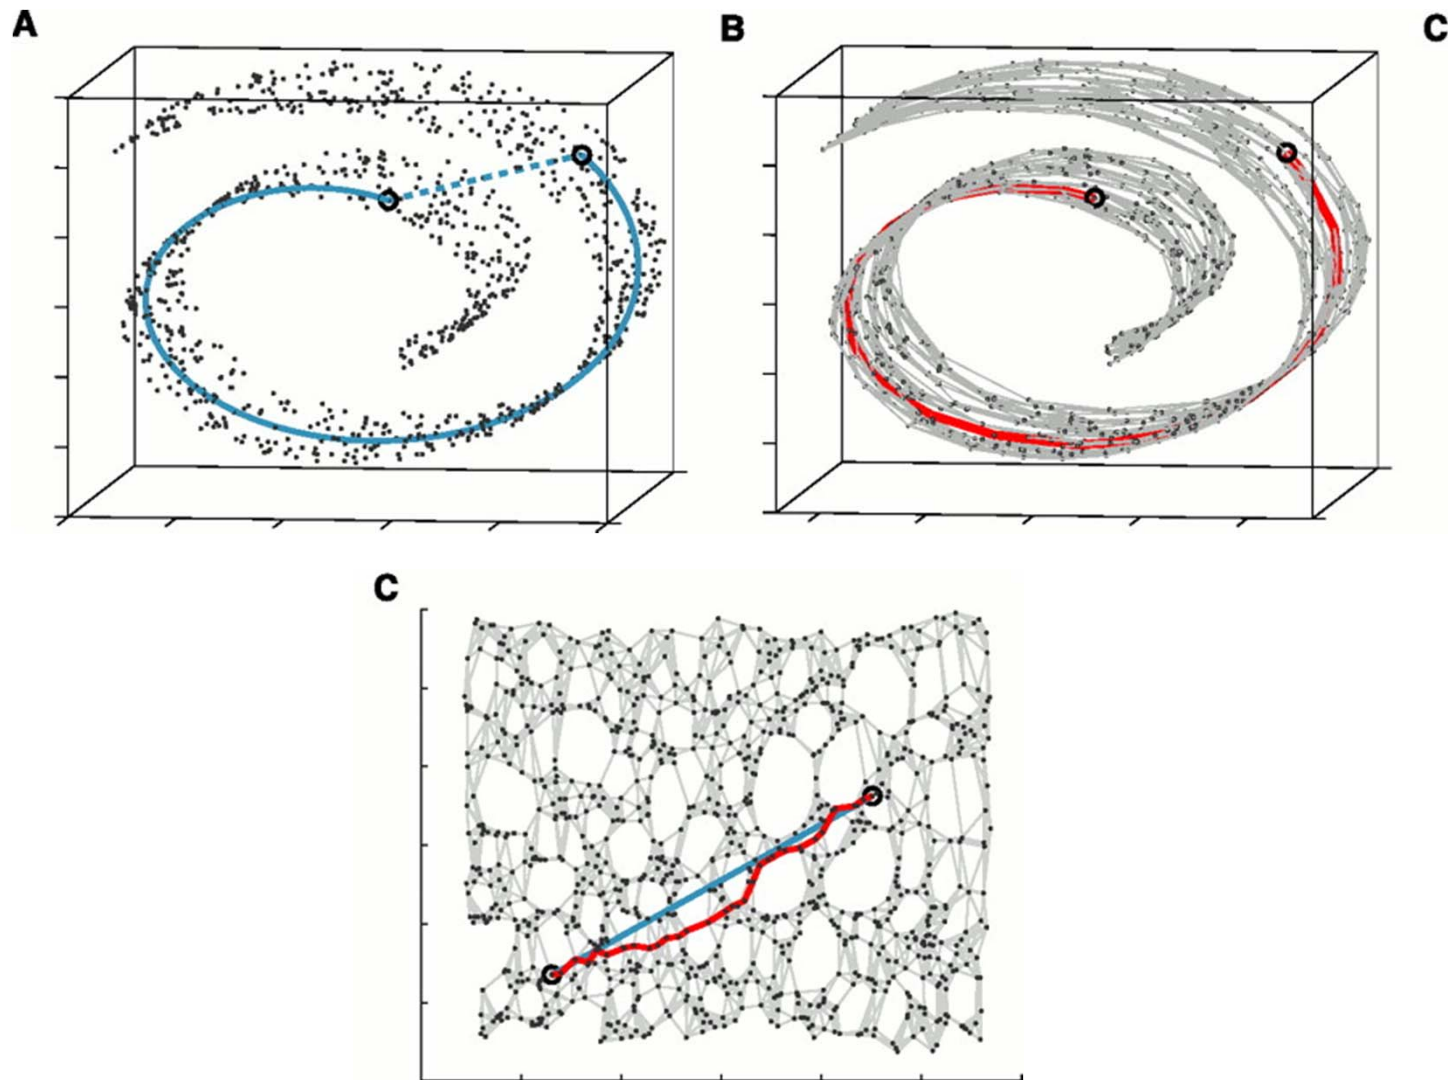
- We can use Multi-Dimensional scaling (MDS), a class of statistical techniques that

Given:

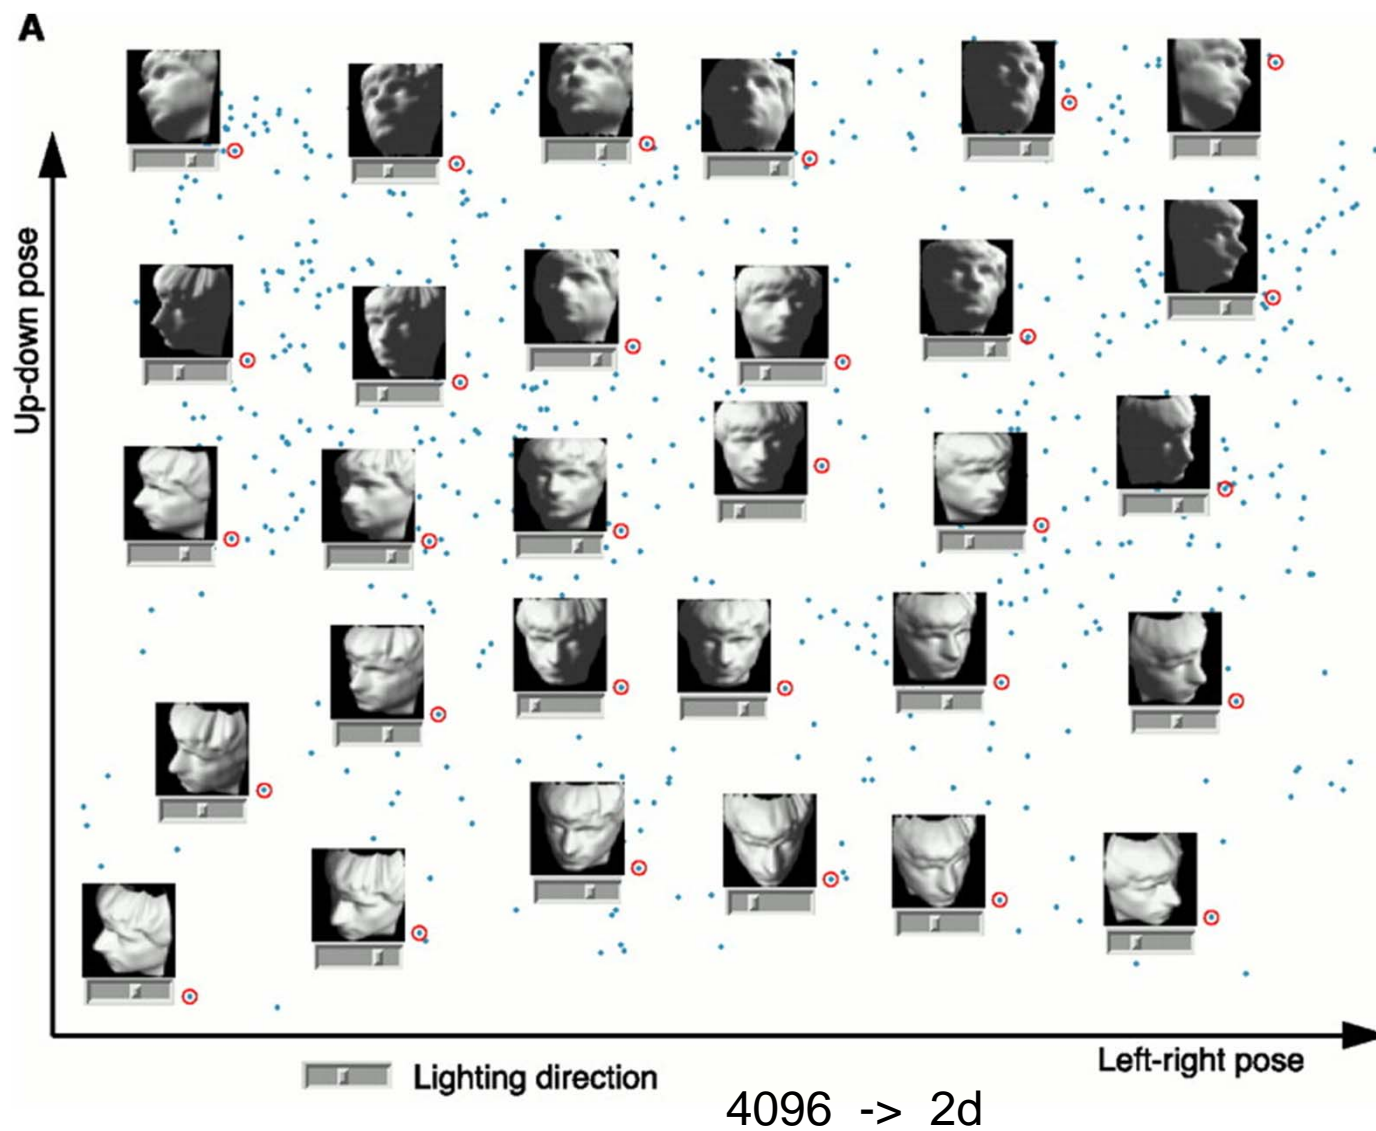
$n \times n$ matrix of dissimilarities between n objects

Outputs: a coordinate configuration of the data in a low-dimensional space R^d whose Euclidean distances closely match given dissimilarities.

ISOMAP on Swiss Roll Data



ISOMAP Examples



ISOMAP Examples

