

Perceptron

A Canonical Representation

- Given a training example: $(\langle x_1, x_2, x_3, x_4 \rangle, y)$ $y \in \{-1, 1\}$
- Transform it to canonical representation

$$(\langle 1, x_1, x_2, x_3, x_4 \rangle, y)$$

- Learn a linear function $g(\mathbf{x}, \mathbf{w}) = \mathbf{w}^T \mathbf{x}$, where $\mathbf{w} = \langle w_0, w_1, w_2, w_3, w_4 \rangle$
- Each \mathbf{w} corresponds to one hypothesis

$$h(\mathbf{x}) = \text{sign}(g(\mathbf{x}, \mathbf{w}))$$

- A prediction is correct if $y \mathbf{w}^T \mathbf{x} > 0$
- Goal of learning is to find a good \mathbf{w}
 - e.g., a \mathbf{w} such that $h(\mathbf{x})$ makes few mis-predictions

Learning w :


An Optimization Problem

- Formulate learning problem as an optimization problems
 - Given:
 - A set of N training examples
 $\{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
 - A loss function L
 - Find the weight vector \mathbf{w} that minimizes the objective function - the expected/average loss on training data

$$J(w) = \frac{1}{N} \sum_{i=1}^N L(w \cdot x_i, y_i)$$

- Many machine learning algorithms apply some optimization algorithm to find a good hypothesis.

Gradient Descent Search




- Start with initial $\mathbf{w} = (w_0, \dots, w_n)$
- Compute gradient $\nabla J(\mathbf{w}_0) = \left(\frac{\partial J(\mathbf{w})}{\partial w_0}, \frac{\partial J(\mathbf{w})}{\partial w_1}, \dots, \frac{\partial J(\mathbf{w})}{\partial w_n} \right)_{\mathbf{w}_0}$
- $w_{t+1} = w_t - \eta \nabla J(w_t)$, Where η is the “step size” parameter
- Repeat until convergence

Remaining question: what objective to use?

Loss Functions

- 0/1 Loss function: $J_{0/1}(w) = \frac{1}{N} \sum_{i=1}^N L(\text{sgn}(w \cdot x_i), y_i)$
 $L(y', y) = 0$ when $y' = y$, otherwise $L(y', y) = 1$
- Does not produce useful gradient since the surface of J is flat




0/1 loss

Loss Functions


- Instead we will consider the “**perceptron criterion**” (a slightly modified version of hinge loss):

$$J_p(w) = \frac{1}{N} \sum_{i=1}^N \max(0, -y_i w \cdot x_i)$$

- The term $\max(0, -y_i w \cdot x_i)$ is 0 when y_i is predicted correctly otherwise it is equal to the “confidence” in the mis-prediction
- Has a nice gradient leading to the solution region



0/1 loss



Perceptron
criterion

Stochastic Gradient Descent

- The objective function consists of a sum over data points---we can update the parameter after observing each example
- This is referred to as Stochastic gradient descent approach

$$J(\mathbf{w}) = \frac{1}{N} \sum_{i=1}^N \max(0, -y_i \mathbf{w} \cdot \mathbf{x}_i)$$

$$J_i(\mathbf{w}) = \max(0, -y_i \mathbf{w} \cdot \mathbf{x}_i)$$

$$\frac{\partial J_i}{\partial w_j} = \begin{cases} 0 & \text{if } y_i \mathbf{w} \cdot \mathbf{x}_i > 0 \\ -y_i x_{ij} & \text{otherwise} \end{cases}$$

$$\nabla J_i = \begin{cases} 0 & \text{if } y_i \mathbf{w} \cdot \mathbf{x}_i > 0 \\ -y_i \mathbf{x}_i & \text{otherwise} \end{cases}$$

After observing (\mathbf{x}_i, y_i) , if it is a mistake $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$

Online Perceptron Algorithm

Let $\mathbf{w} \leftarrow (0,0,0,\dots,0)$

Repeat

Accept training example $i : (\mathbf{x}_i, y_i)$

$u_i \leftarrow \mathbf{w} \cdot \mathbf{x}_i$

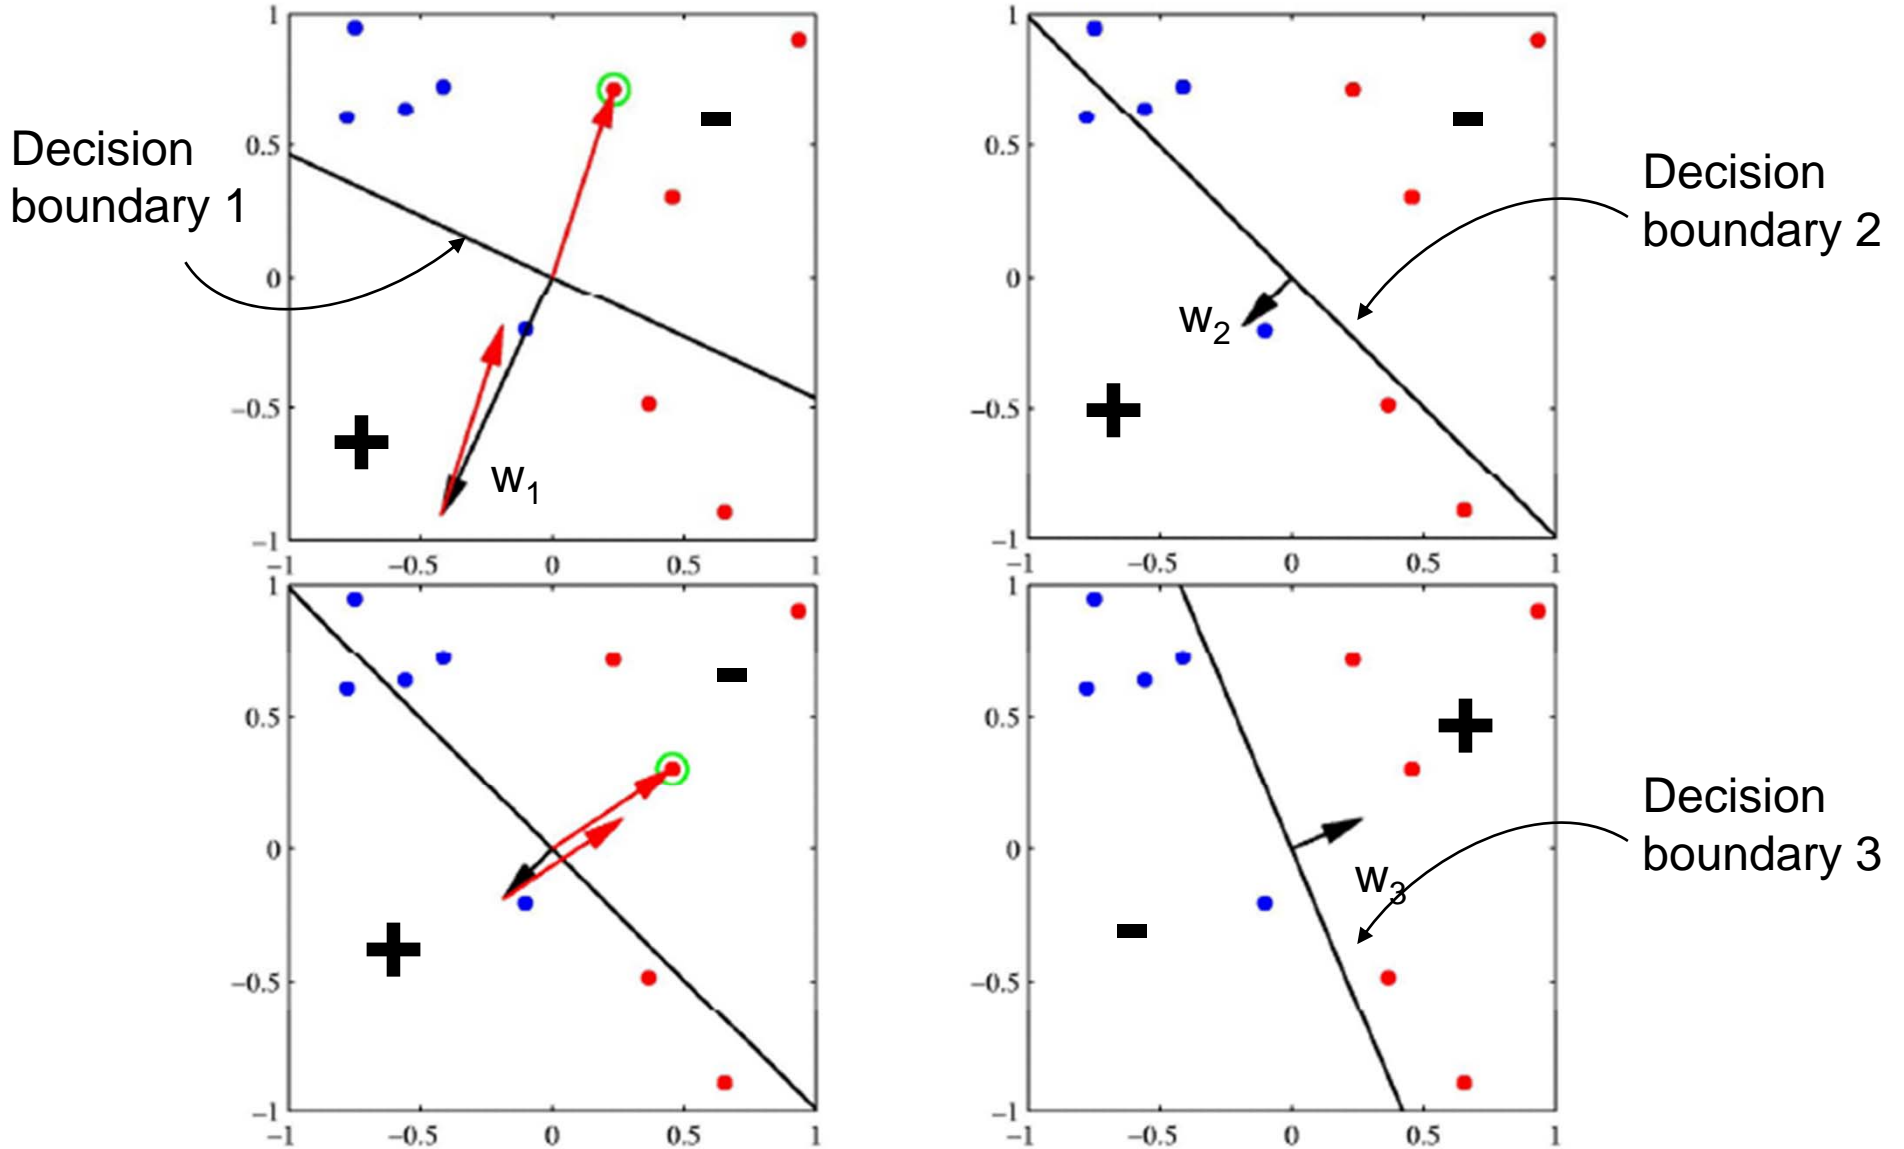
if $y_i \cdot u_i \leq 0$

$\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$

Online learning refers to the learning mode in which the model update is performed each time a single observation is received.

Batch learning in contrast performs model update after observing the whole training set.

When an error is made, moves the weight in a direction that corrects the error



Red points belong to the positive class,
blue points belong to the negative class

Batch Perceptron Algorithm

Given : training examples (\mathbf{x}_i, y_i) , $i = 1, \dots, N$

Let $\mathbf{w} \leftarrow (0, 0, 0, \dots, 0)$

do

$\mathit{delta} \leftarrow (0, 0, 0, \dots, 0)$

for $i = 1$ to N do

$u_i \leftarrow \mathbf{w} \cdot \mathbf{x}_i$

if $y_i \cdot u_i \leq 0$

$\mathit{delta} \leftarrow \mathit{delta} - y_i \cdot x_i$

$\mathit{delta} \leftarrow \mathit{delta} / N$

$\mathbf{w} \leftarrow \mathbf{w} - \eta \mathit{delta}$

until $|\mathit{delta}| < \varepsilon$

Simplest case: $\eta = 1$ and don't normalize – 'Fixed increment perceptron'

η – the step size

- Also referred as the learning rate
- In practice, recommend to decrease η as learning continues
- Some optimization approaches set step-size automatically, e.g., by line search, and converge faster
- If linearly separable, there is only one basin for the hinge loss, thus local minimum is the global minimum

Online VS. Batch Perceptron

- Batch learning learns with a batch of examples collectively
- Online learning learns with one example at a time
- Both learning mechanisms are useful in practice
- Online Perceptron is sensitive to the order that training examples are received
- In batch training, the correction incurred by each mistake is accumulated and applied at once at the end of the iteration
- In online training, each correction is applied immediately once a mistake is encountered, which will change the decision boundary, thus different mistakes maybe encountered for online and batch training
- Online training performs stochastic gradient descent, an approximation to real gradient descent, which is used by the batch training

Convergence Theorem

(Block, 1962, Novikoff, 1962)

Given training example sequence $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)$.

If $\forall i, \|\mathbf{x}_i\| \leq D$, and $\exists \mathbf{u}, \|\mathbf{u}\| = 1$ and $y_i \mathbf{u} \cdot \mathbf{x}_i \geq \gamma > 0$ for all i ,

then the number of mistakes that the perceptron algorithm makes is at most $(D / \gamma)^2$.

Note that $\|\cdot\|$ is the Euclidean length of a vector.

Proof

To show convergence, we just need to show that each update moves the weight vector closer to a solution vector by a lower bounded amount

Note that $\alpha \mathbf{u}$ is also a solution vector, given that \mathbf{u} is a solution vector, where α is an arbitrary scaling factor

Let \mathbf{x}_k be the k th mistake, we have $\mathbf{w}(k+1) = \mathbf{w}(k) + y_k \mathbf{x}_k$

$$\begin{aligned} & \|\mathbf{w}(k+1) - \alpha \mathbf{u}\|^2 \\ &= \|\mathbf{w}(k) + y_k \mathbf{x}_k - \alpha \mathbf{u}\|^2 = \|(\mathbf{w}(k) - \alpha \mathbf{u}) + y_k \mathbf{x}_k\|^2 \\ &= \|\mathbf{w}(k) - \alpha \mathbf{u}\|^2 + 2y_k [\mathbf{x}_k^T (\mathbf{w}(k) - \alpha \mathbf{u})] + (y_k)^2 \|\mathbf{x}_k\|^2 \\ &= \|\mathbf{w}(k) - \alpha \mathbf{u}\|^2 + 2y_k \mathbf{x}_k^T \mathbf{w}(k) - 2y_k \alpha \mathbf{u}^T \mathbf{x}_k + \|\mathbf{x}_k\|^2 \\ &\leq \|\mathbf{w}(k) - \alpha \mathbf{u}\|^2 + 2y_k \mathbf{x}_k^T \cdot \mathbf{w}(k) - 2y_k \alpha \mathbf{u}^T \mathbf{x}_k + D^2, \text{ because } \|\mathbf{x}_k\| \leq D \\ &\leq \|\mathbf{w}(k) - \alpha \mathbf{u}\|^2 - 2\alpha \mathbf{u}^T \mathbf{x}_k + D^2, \text{ because } y_k \mathbf{x}_k^T \mathbf{w}(k) \leq 0 \\ &\leq \|\mathbf{w}(k) - \alpha \mathbf{u}\|^2 - 2\alpha \gamma + D^2, \text{ because } y_k \mathbf{u}^T \mathbf{x}_k \geq \gamma \end{aligned}$$

Because α is an arbitrary scaling factor, we can set $\alpha = \frac{D^2}{\gamma}$

$$\|\mathbf{w}(k+1) - \alpha \mathbf{u}\|^2 \leq \|\mathbf{w}(k) - \alpha \mathbf{u}\|^2 - D^2$$

Proof (cont.)

By induction on k , we can show that

$$\|\mathbf{w}(k+1) - \alpha \mathbf{u}\|^2 \leq \|\mathbf{w}(1) - \alpha \mathbf{u}\|^2 - kD^2 = \alpha^2 \|\mathbf{u}\|^2 - kD^2 = \alpha^2 - kD^2$$

$$\Leftrightarrow \alpha^2 - kD^2 \geq 0$$


$$\Leftrightarrow k \leq \frac{\alpha^2}{D^2} \quad \left(\alpha = \frac{D^2}{\gamma}\right)$$

$$\Leftrightarrow k \leq (D/\gamma)^2$$

Margin

- γ is referred to as the margin
 - The minimum distance from data points to the decision boundary
 - The bigger the margin, the easier the classification problem is
 - The bigger the margin, the more confident we are about our prediction
- We will see later in the course this concept leads to one of the recent most exciting developments in the ML field – support vector machines

Not linearly separable case



- In such cases the algorithm will never stop! How to fix?
- Look for decision boundary that make as few mistakes as possible – NP-hard!

Fixing the Perceptron

- Idea one: only go through the data once, or a fixed number of times

```
Let  $\mathbf{w} \leftarrow (0,0,0,...,0)$   
Repeat for N times  
    Take a training example  $i : (\mathbf{x}_i, y_i)$   
     $u_i \leftarrow \mathbf{w} \cdot \mathbf{x}_i$   
    if  $y_i \cdot u_i \leq 0$   
         $\mathbf{w} \leftarrow \mathbf{w} + y_i \mathbf{x}_i$ 
```

- At least this stops
- Problem: the final \mathbf{w} might not be good e.g. right before it stops the algorithm might perform an update on a total outlier

Voted Perceptron

- Keep intermediate hypotheses and have them vote [Freund and Schapire 1998]

```
Let  $w_0 = (0,0,0, \dots, 0)$   
 $c_0 = 0$   
repeat  
    Take example  $i : (x_i, y_i)$   
     $u_i \leftarrow \mathbf{w}_n \cdot \mathbf{x}_i$   
    if  $y_i \cdot u_i \leq 0$   
         $\mathbf{w}_{n+1} \leftarrow \mathbf{w}_n + y_i \mathbf{x}_i$   
         $c_{n+1} = 0$   
         $n = n + 1$   
    else  
         $c_n = c_n + 1$ 
```

Store a collection of linear separators $\mathbf{w}_0 \mathbf{w}_1 \dots$, along with their survival time $c_0, c_1 \dots$

The c 's can be good measures of the reliability of the \mathbf{w} 's

For classification, take a weighted vote among all separators:

$$\text{sgn}\left\{\sum_{n=0}^N c_n \text{sgn}(\mathbf{w}_n^T \mathbf{x})\right\}$$

Summary of Perceptron

- Learns a Classifier $\hat{y} = f(\mathbf{x})$ directly
- Applies gradient descent search to optimize the hinge loss function
 - Online version performs stochastic gradient descent
- Guaranteed to converge in finite steps if linearly separable
 - There exists an upper bound on the number of corrections needed
 - Inversely proportional to the margin of the optimal decision boundary
- If not linearly separable, use voted perceptrons


Geometric Interpretation of Linear Discriminant Functions

- Two classes

$$y(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + w_0$$

if $y(\mathbf{x}) \geq 0$, assign to C_1
otherwise, assign to C_2

- Decision boundary: $y(\mathbf{x})=0$
- Decision boundary is perpendicular to \mathbf{w}



The signed distance (positive if \mathbf{x} is on the positive side, negative otherwise) from any point \mathbf{x} to the decision boundary is: $\frac{y(\mathbf{x})}{\|\mathbf{w}\|}$

Note that in Perceptron, due to the adoption of the canonical representation, all training points will lie on the hyperplane $x_0=0$, and the decision boundary will always go through the origin.